
DOCKER INTRODUCTION



JULIEN MAITREHENRY

Dev & Ops @PetalMD

<https://github.com/jmaitrehenry>

@jmaitrehenry

AGENDA

- ▶ The Challenge
- ▶ Docker as solution
- ▶ But What is Docker?
- ▶ What's an image?
- ▶ What's a container?
- ▶ Questions?

THE CHALLENGE

THE CHALLENGE

Multiplicity of Stacks

 **Static website**
nginx 1.5 + modsecurity + openssl + bootstrap 2

 **Background workers**
Python 3.0 + celery + pyredis + libcurl + ffmpeg + libopencv + nodejs + phantomjs

 **User DB**
postgresql + pgv8 + v8

 **Queue**
Redis + redis-sentinel

 **Analytics DB**
hadoop + hive + thrift + OpenJDK

 **Web frontend**
Ruby + Rails + sass + Unicorn

 **API endpoint**
Python 2.7 + Flask + pyredis + celery + psycopg + postgresql-client

Do services and apps
interact
appropriately?

Multiplicity of
hardware
environments

 **Development VM**

 **QA server**

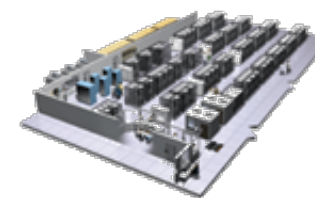
Customer Data Center



Public Cloud

Disaster recovery

Production Servers



Production Cluster




Contributor's laptop



Can I migrate
smoothly and
quickly?

THE CHALLENGE

THE MATRIX FROM HELL

	Static website	?	?	?	?	?	?	?
	Web frontend	?	?	?	?	?	?	?
	Background workers	?	?	?	?	?	?	?
	User DB	?	?	?	?	?	?	?
	Analytics DB	?	?	?	?	?	?	?
	Queue	?	?	?	?	?	?	?
		Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers
								

THE CHALLENGE

CARGO TRANSPORT PRE-1960

Multiplicity of Goods



Do I worry about
how goods interact
(e.g. coffee beans
next to spices)








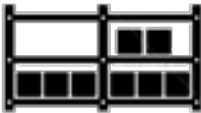





Multiplicity of
methods for
transporting/storing



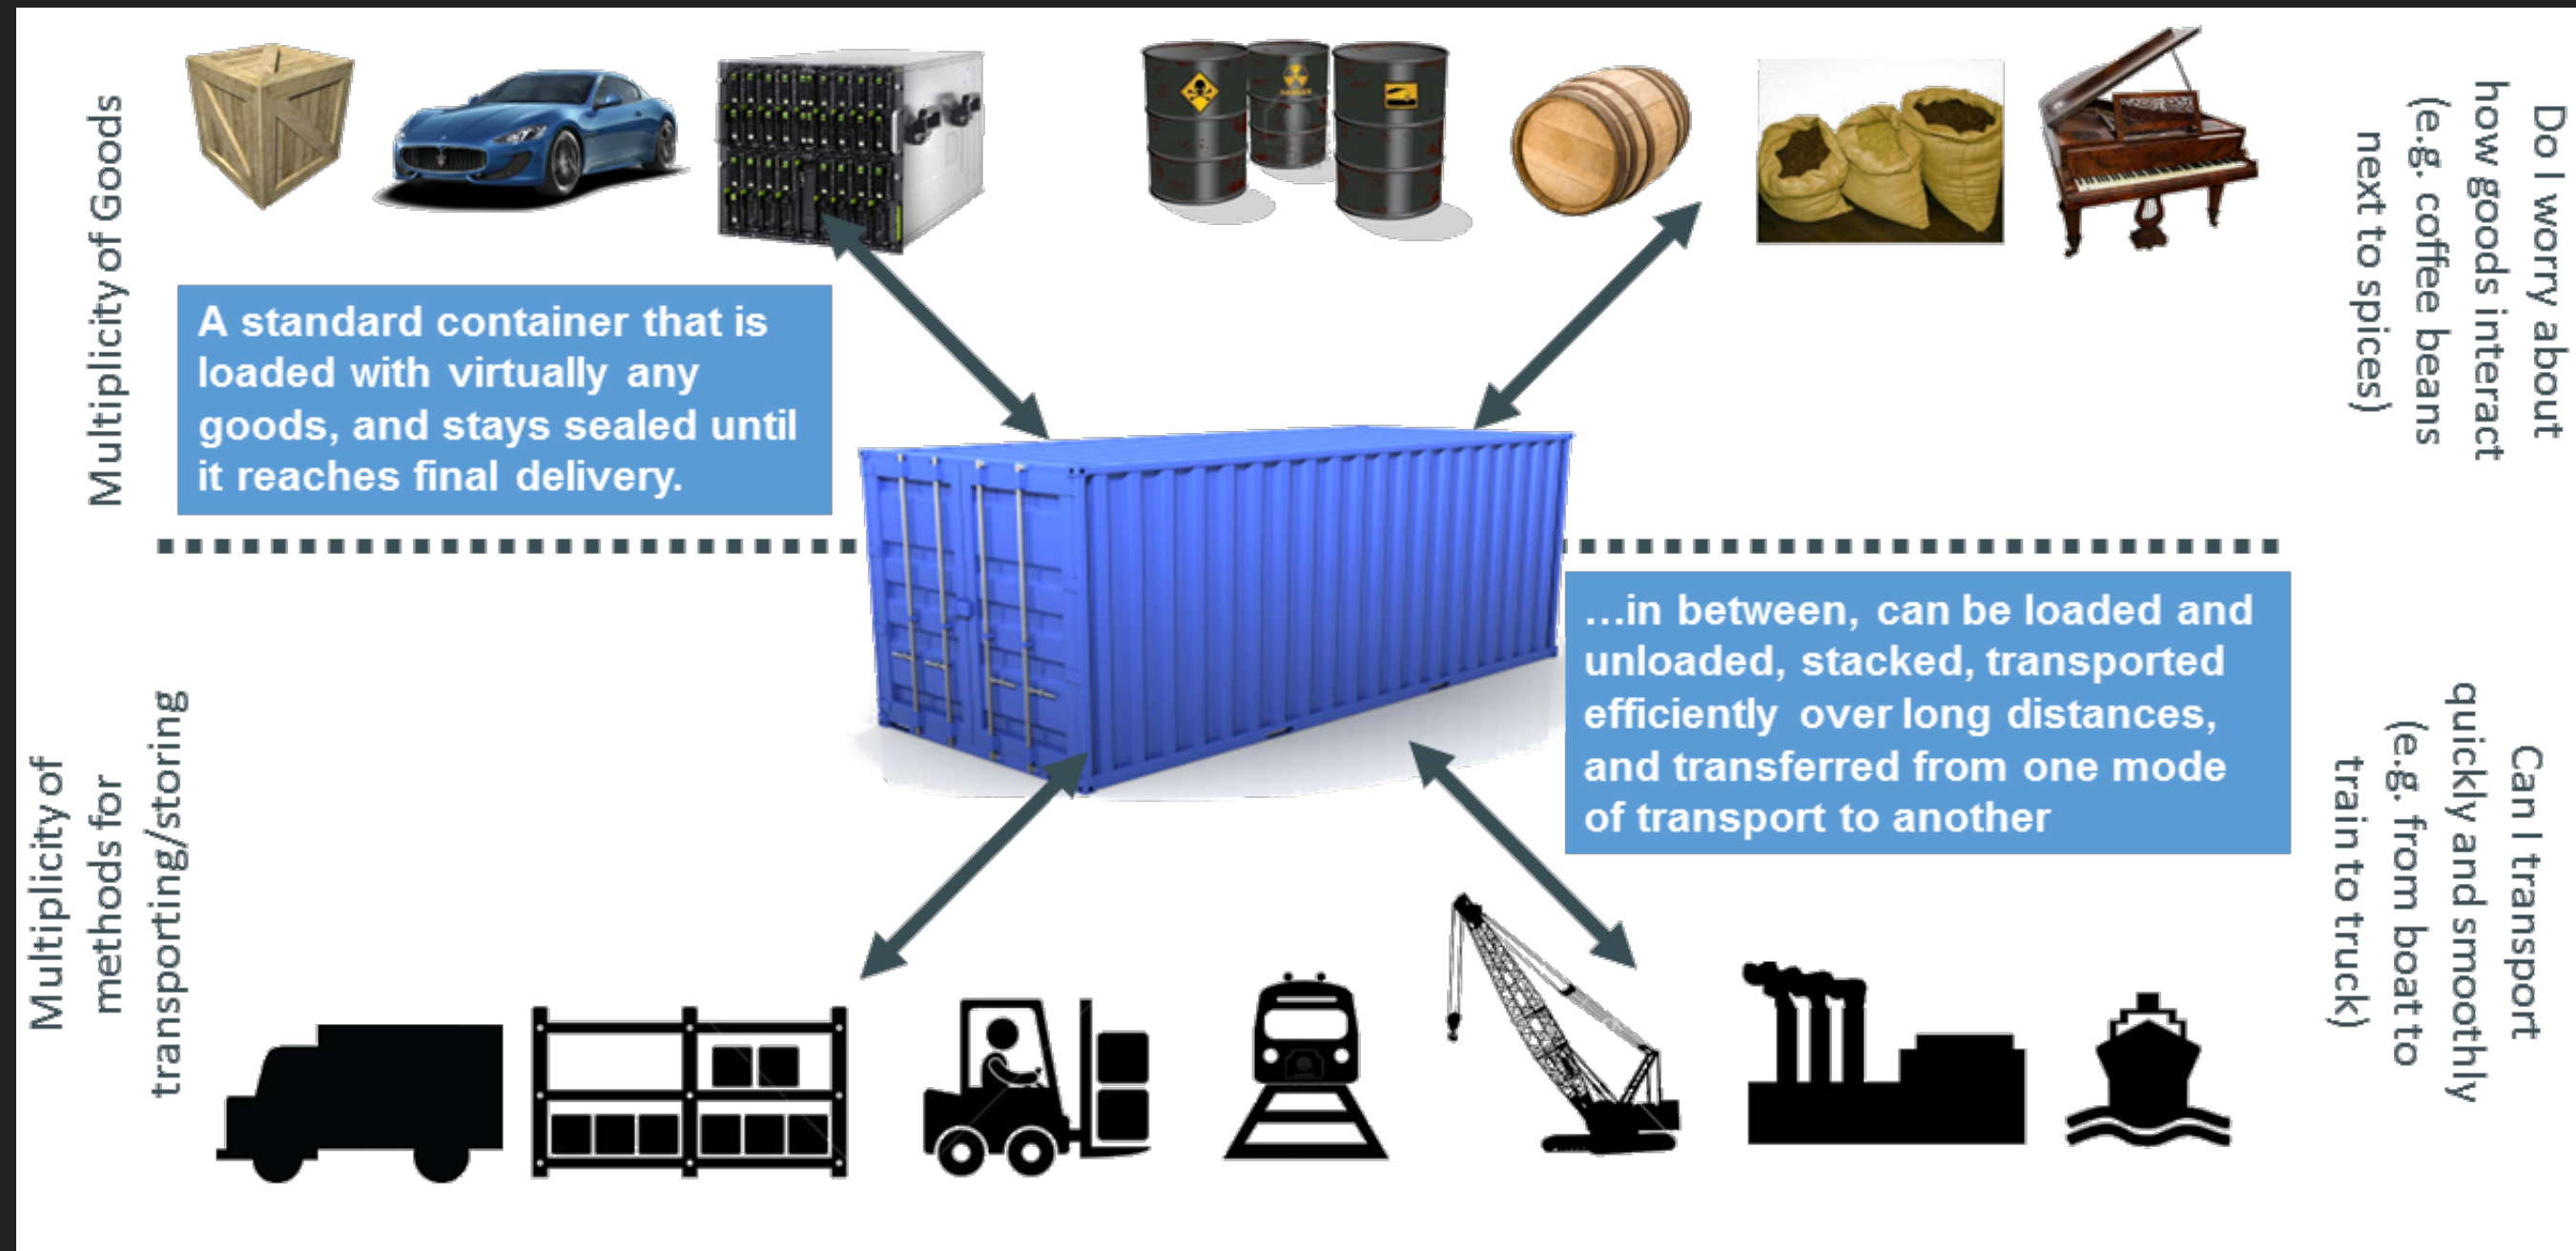
Can I transport quickly
and smoothly
(e.g. from boat to train
to truck)

THE CHALLENGE

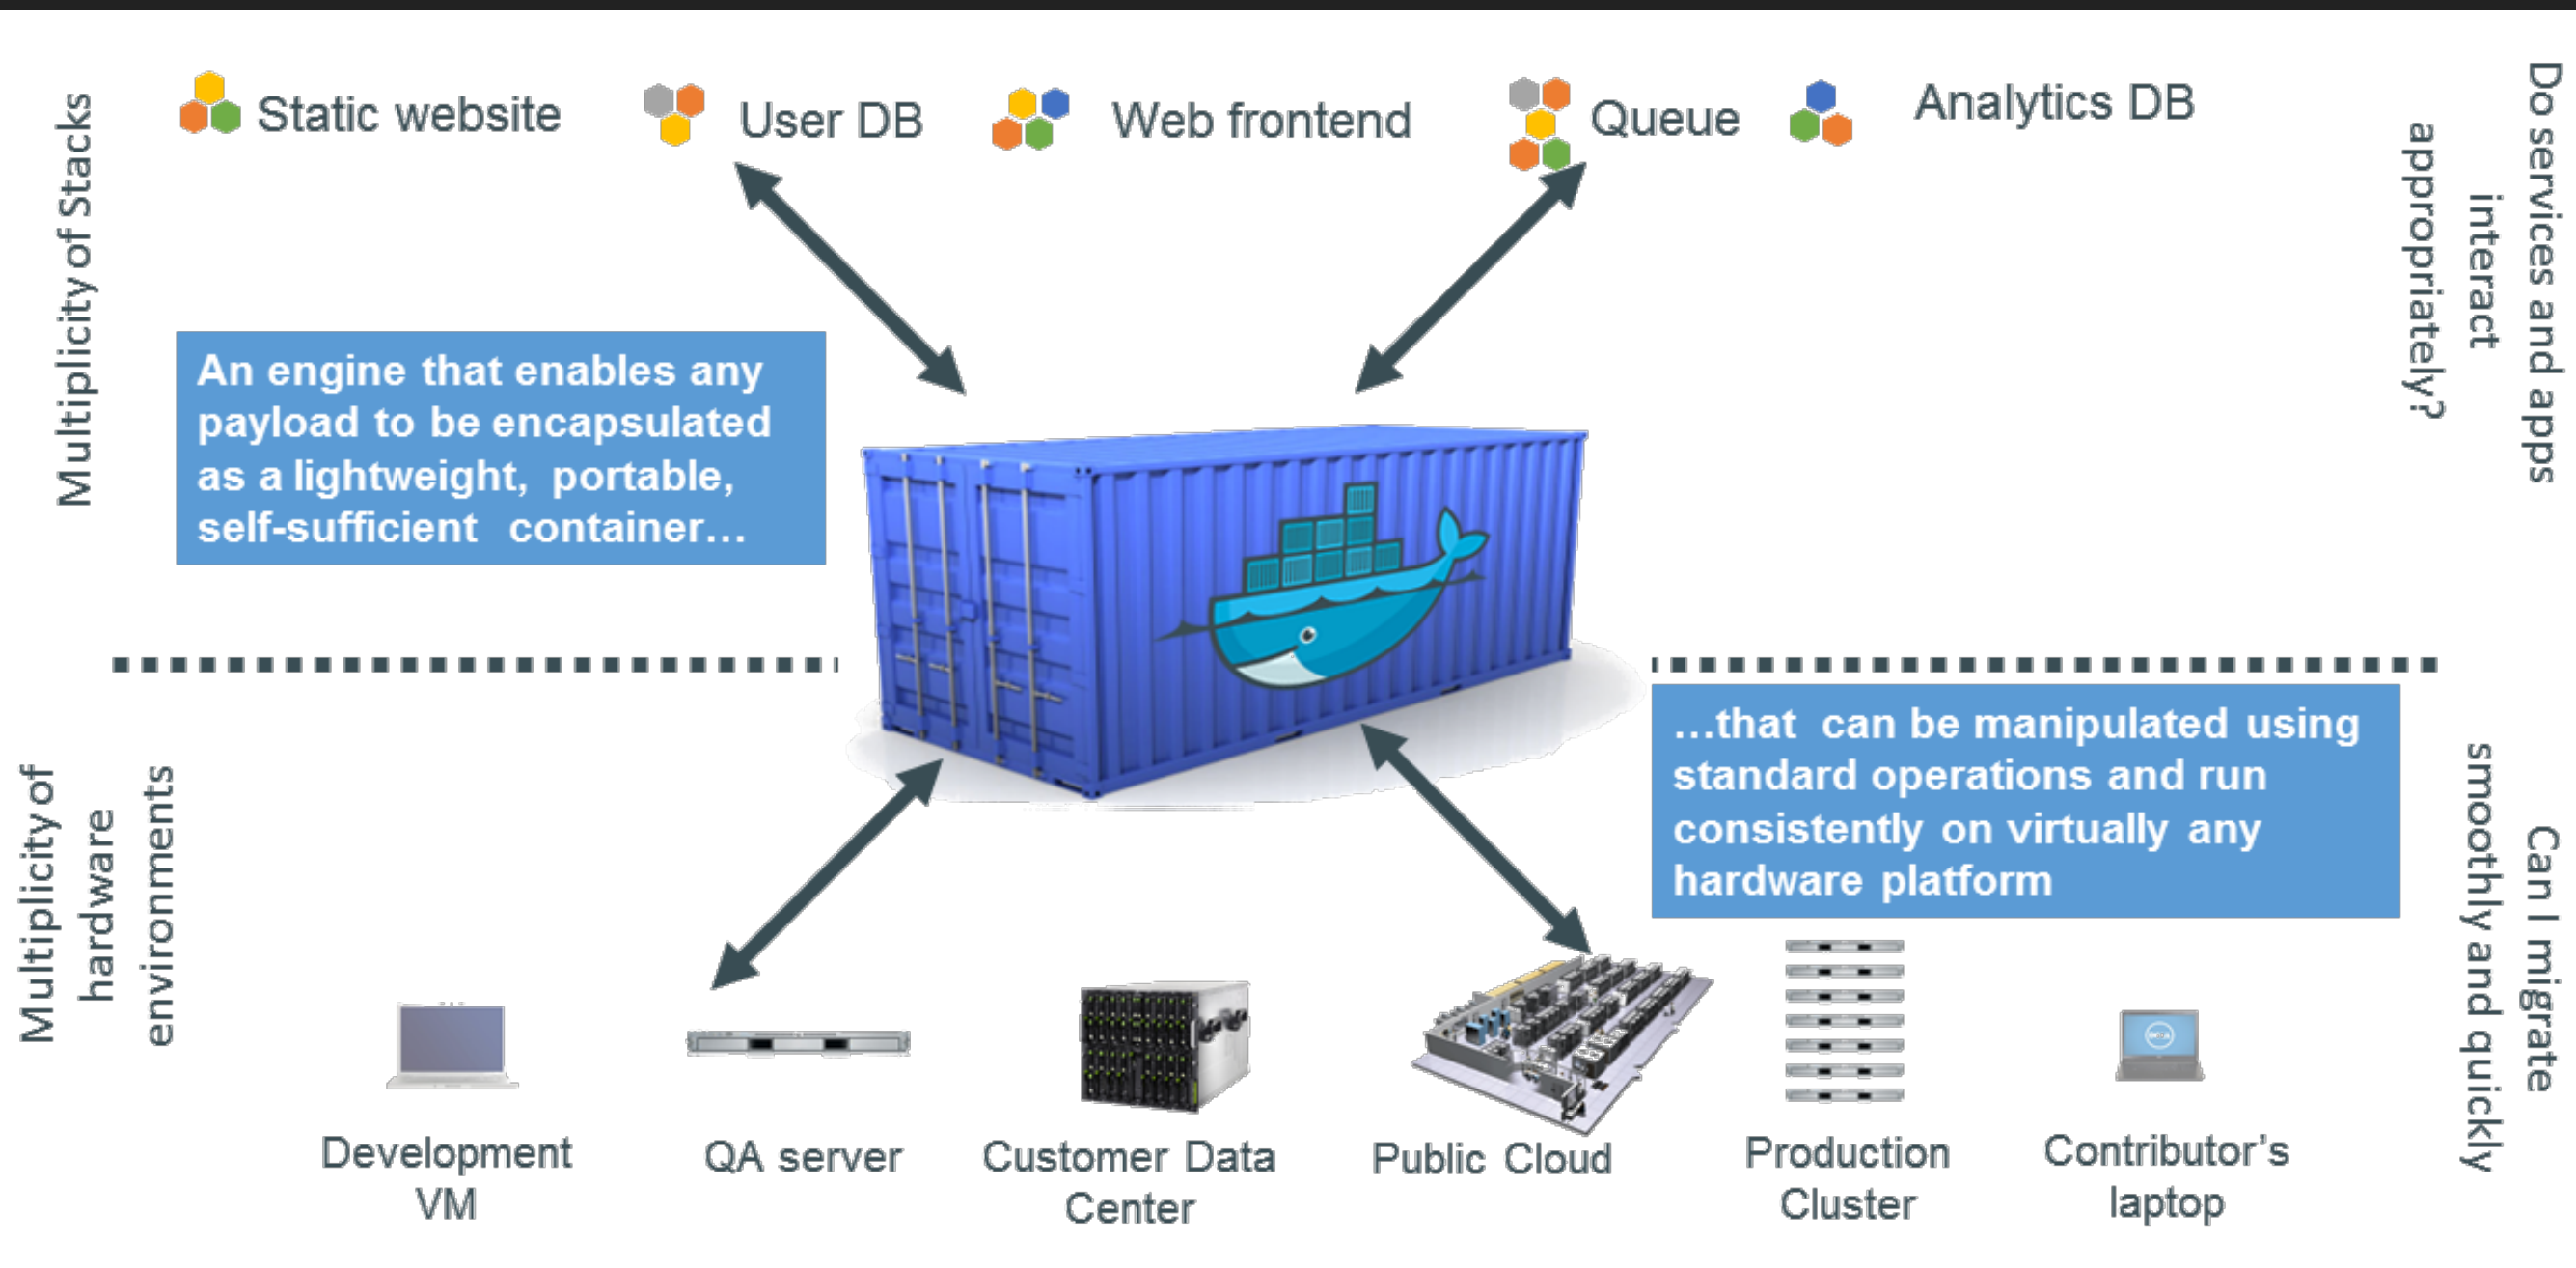
ALSO A MATRIX FROM HELL

	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
							

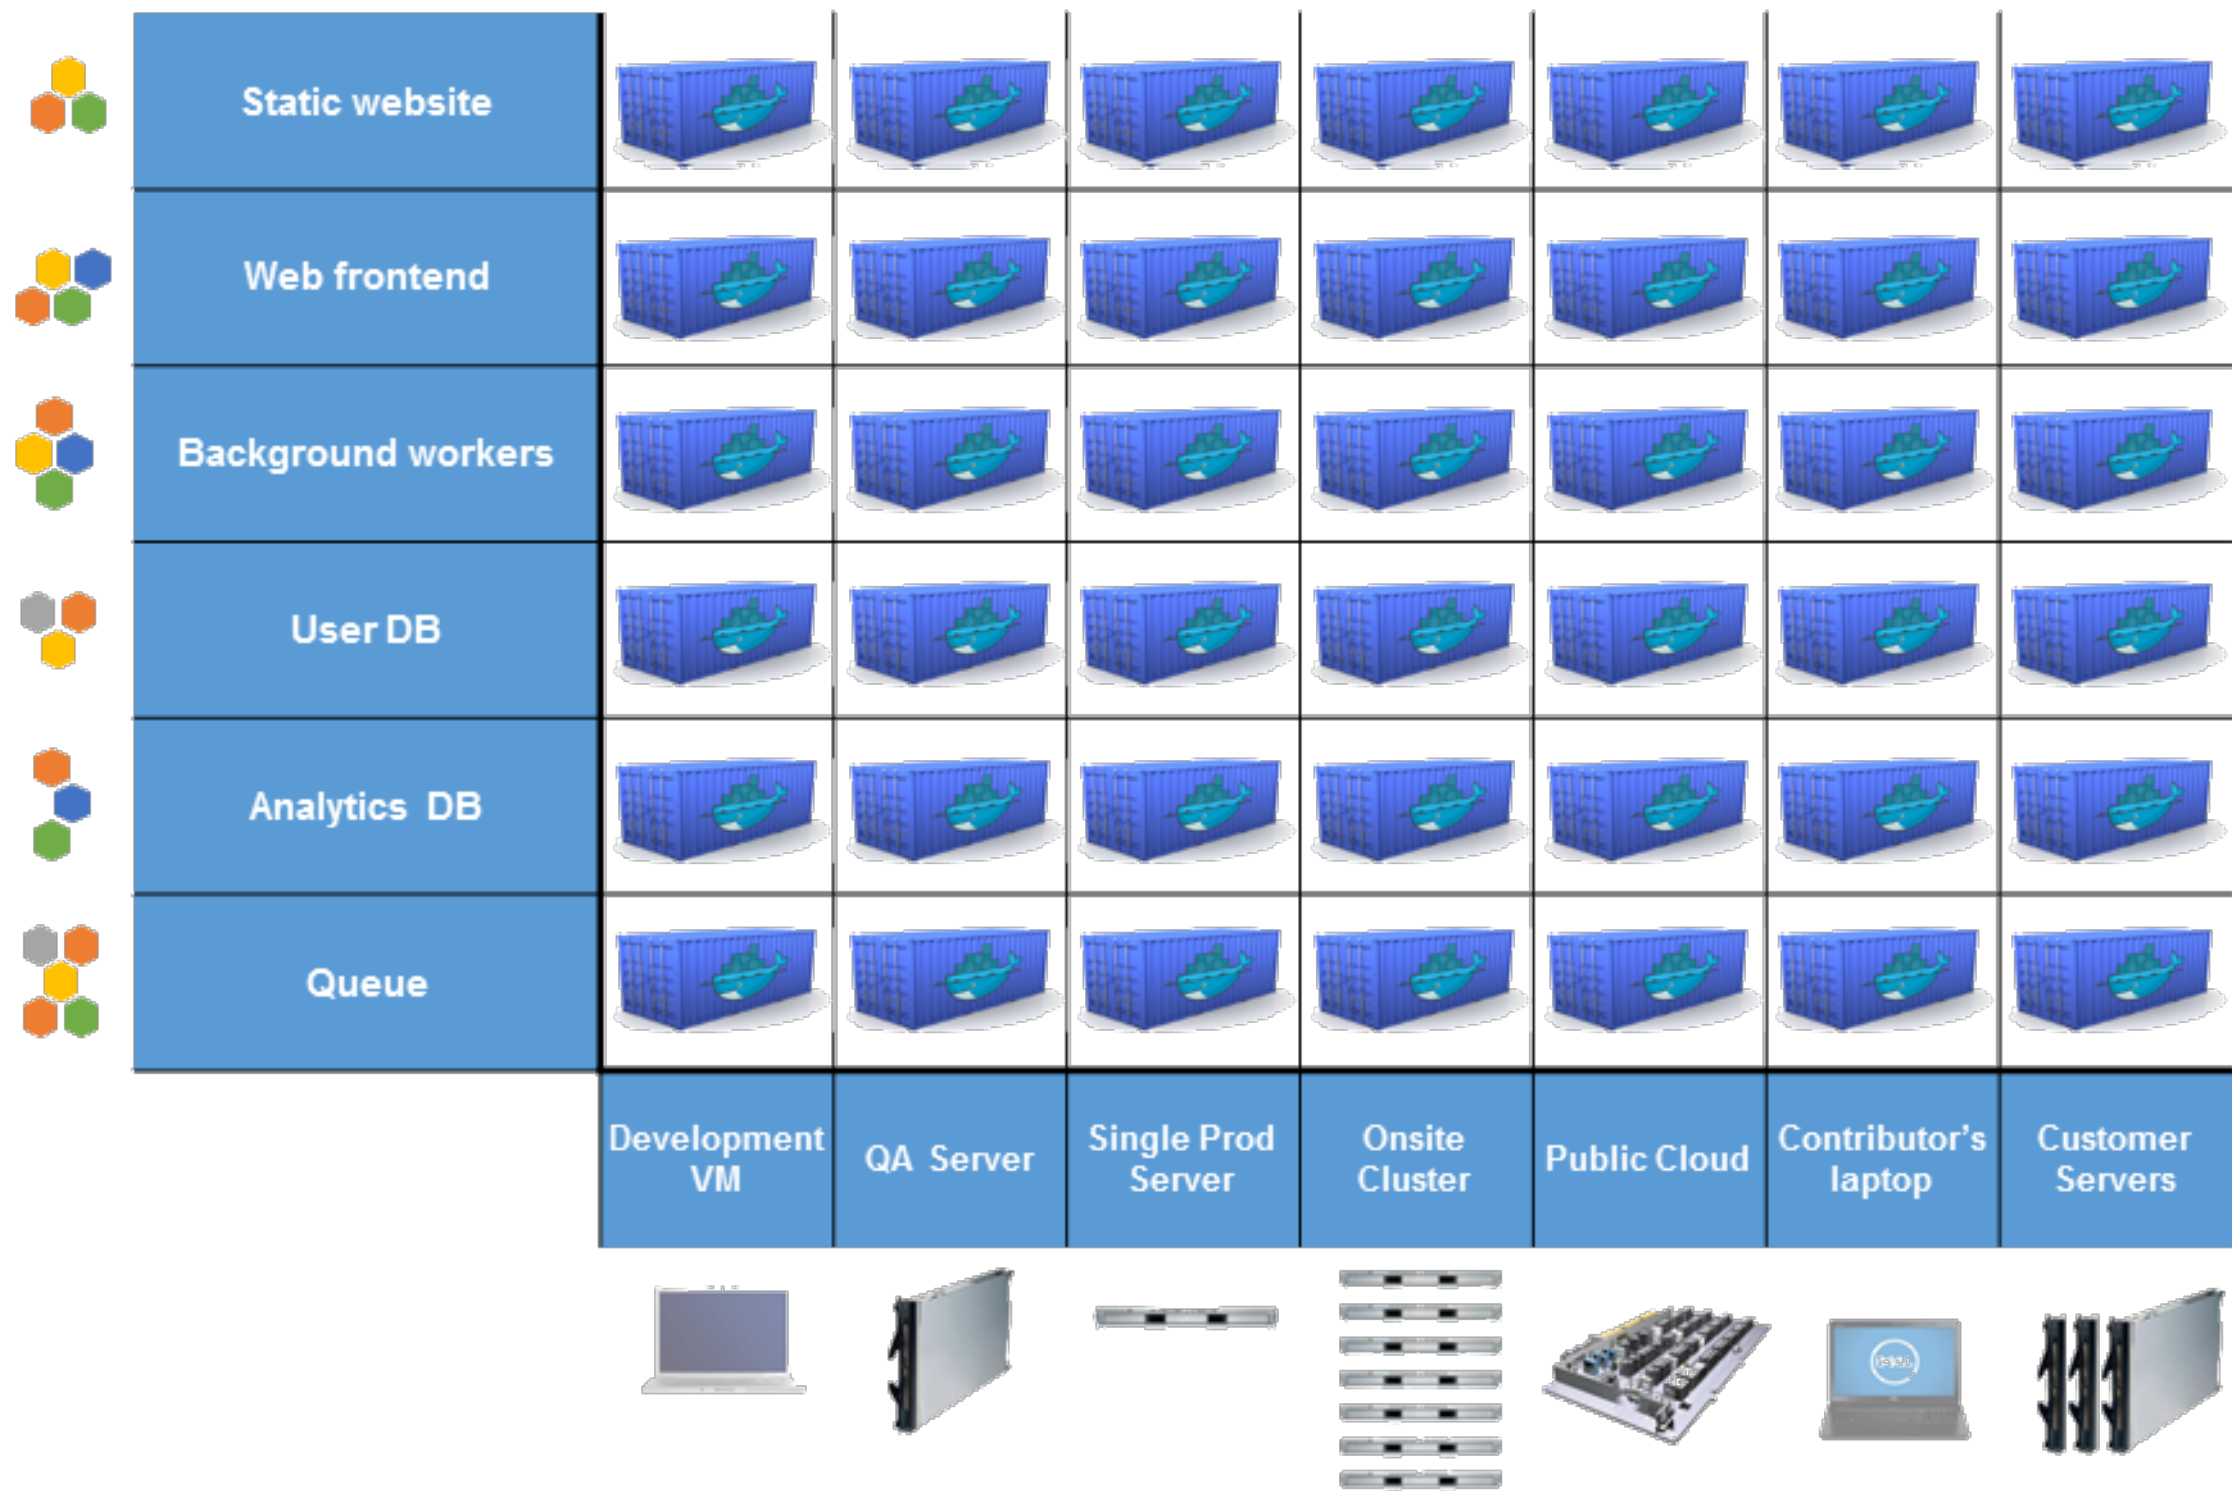
SOLUTION: INTERMODAL SHIPPING CONTAINER



DOCKER IS A CONTAINER SYSTEM FOR CODE



DOCKER ELIMINATES THE MATRIX FROM HELL



**BUT WHAT IS
DOCKER?**

Docker containers wrap a piece of software in a complete filesystem that contains everything needed to run: code, runtime, system tools, system libraries — anything that can be installed on a server. This guarantees that the software will always run the same, regardless of its environment

BUT WHAT IT'S DOCKER?

DOCKER STATS



32,000+
GitHub Stars



4B+
Docker Container Downloads



450,000+
Dockerized Apps In Docker Hub



250+
Meetup Groups In 70+ Countries



2900+
Community Contributors

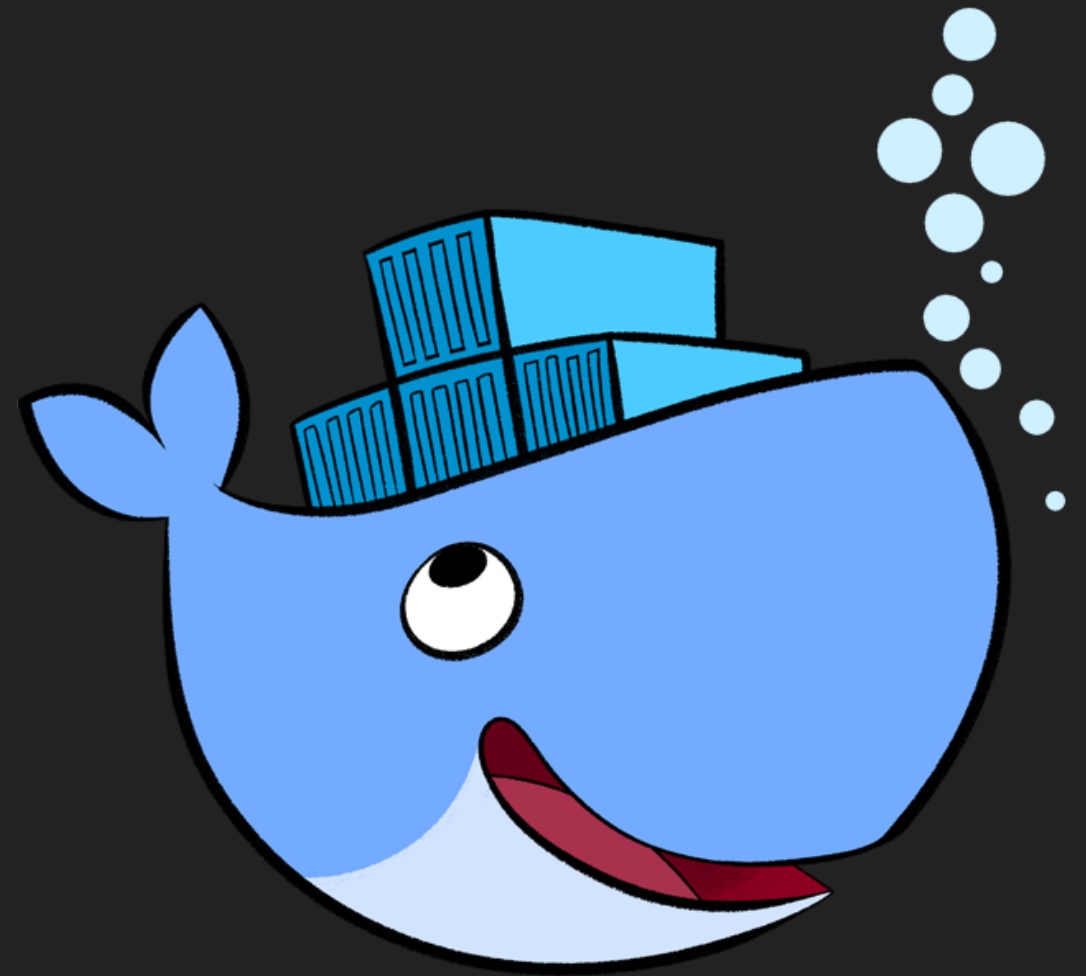


95,000+
Third Party Projects Using Docker

BUT WHAT IT'S DOCKER?

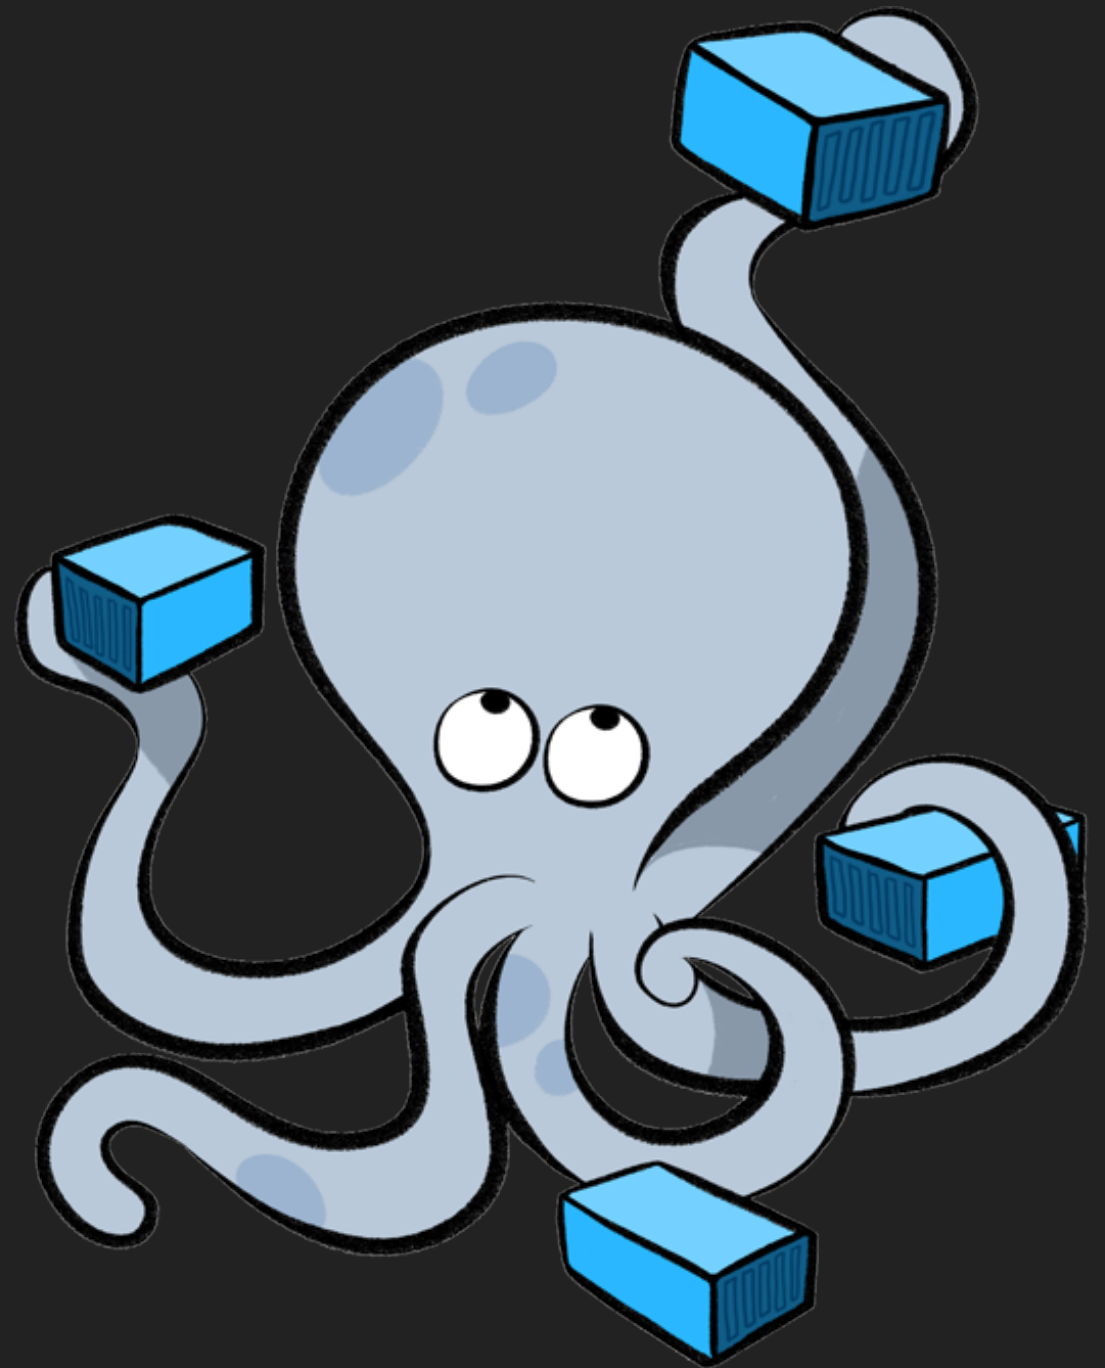
DOCKER ENGINE

- ▶ Lightweight runtime and robust tooling that:
 - ▶ Build image
 - ▶ Ship image
 - ▶ Run containers



DOCKER COMPOSE

- ▶ Define multi-container app with multiple dependencies in one file
- ▶ Hold structure and configuration in a single place
- ▶ Spin application up (or down) with a single command



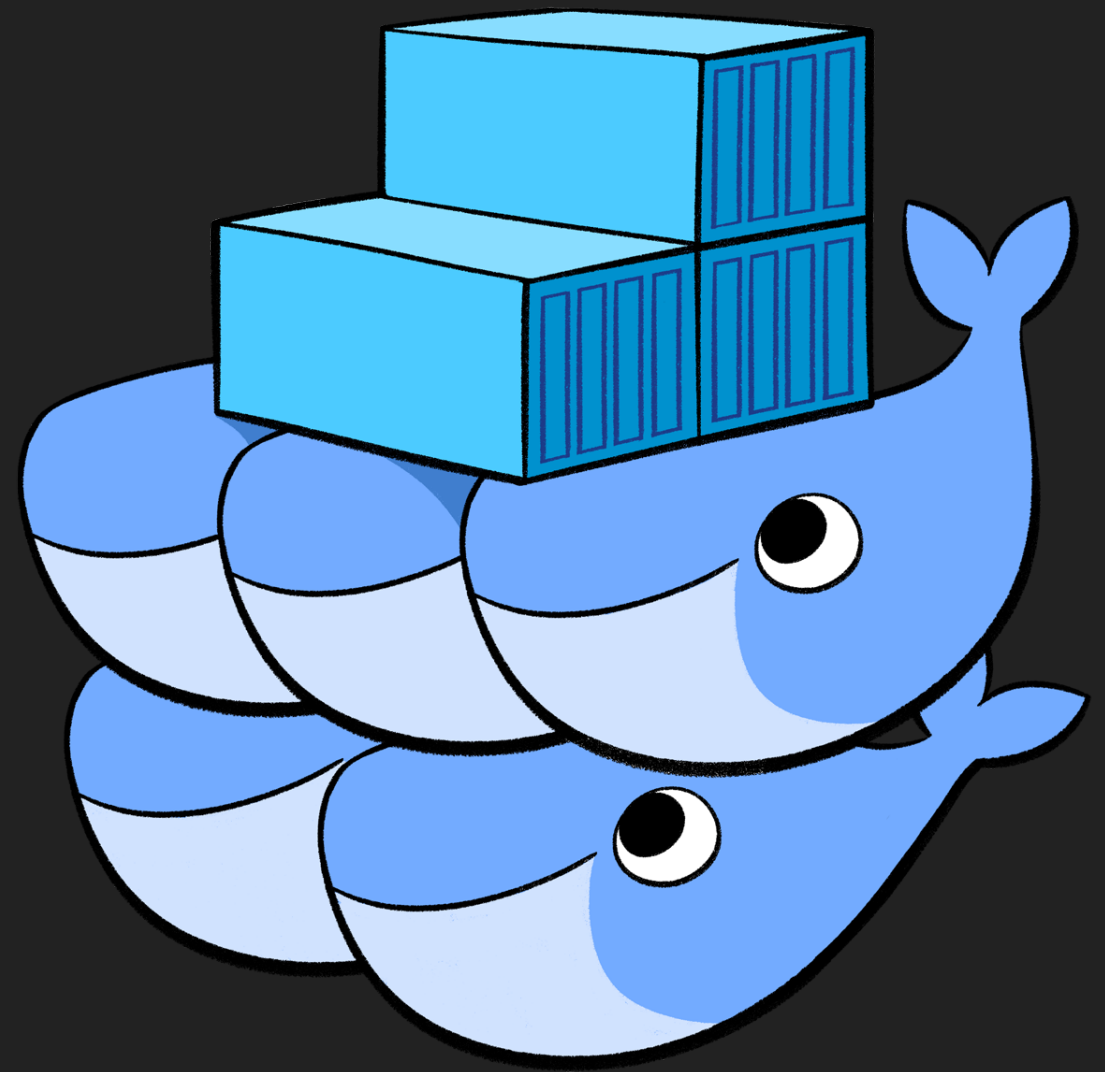
DOCKER REGISTRY

- ▶ Store and distribute Docker Image
- ▶ Many hosted registries available:
 - ▶ Docker Hub, AWS ECR, ...
- ▶ Could be self hosted with many storage backends:
 - ▶ Local, AWS S3, Ceph, OpenStack Swift, ...



DOCKER SWARM

- ▶ Turn group of Docker Engine into a single, virtual Docker Engine
- ▶ Native clustering capabilities
 - ▶ High Scalability
 - ▶ Failover & High Availability
 - ▶ Built-in scheduler
 - ▶ Node discovery
 - ▶ ...



SEPARATION OF CONCERNS: FRANCIS THE DEVELOPER

- ▶ Inside my container:
 - ▶ My code
 - ▶ My libraries
 - ▶ My package manager
 - ▶ My app
 - ▶ My data

SEPARATION OF CONCERNS: JULIEN THE OPS GUY

- ▶ Outside the container:
 - ▶ Logging
 - ▶ Remote access
 - ▶ Network configuration
 - ▶ Monitoring
 - ▶ High Availability

**WHAT'S AN
IMAGE?**

DEFINITION

- ▶ Each Docker image references a list of read-only layers that represent filesystem differences
- ▶ Layers are stacked on top of each other to form a base for a container's root filesystem
- ▶ Each layer are Read Only
- ▶ The Docker storage driver is responsible for stacking these layers and providing a single unified view
- ▶ Each layer can be shared with other image

IT'S A SIMPLE FILE – A DOCKERFILE

```
FROM mhart/alpine-node:6

ENV SERVICE_3000_NAME selfpro
WORKDIR /src

ADD .
RUN apk add --update --no-cache git && \
    npm install

EXPOSE 3000

CMD npm start
```

BUILDING...

```
docker build -t petalmd/selfpro -f docker/Dockerfile.app .
```

```
Sending build context to Docker daemon 1.398 MB
Step 1 : FROM mhart/alpine-node:6
----> 2e8721f40082
Step 2 : ENV SERVICE 3000 NAME selfpro
----> Running in cd8e6156fbec
----> 9d9b1f86a696
Removing intermediate container cd8e6156fbec

[...]

Step 7 : CMD npm start
----> Running in 71bab8276b39
----> cce4c583c3a0
Removing intermediate container 71bab8276b39
Successfully built cce4c583c3a0
```


RESULT ON A STACK OF LAYERS

docker history cce4c583c3a0
docker history petalmd/selfpro

IMAGE	CREATED	CREATED BY	SIZE
cce4c583c3a0	2 minutes ago	/bin/sh -c #(nop) CMD ["/bin/sh" "-c" "npm s	0 B
110c49285fe1	2 minutes ago	/bin/sh -c #(nop) EXPOSE 3000/tcp	0 B
579469b24848	2 minutes ago	/bin/sh -c apk add --update --no-cache git &&	54.16 MB
5b34c5ef9a7f	3 minutes ago	/bin/sh -c #(nop) %s %s in %s ADD dir:b07dd1	1.044 MB
6ba4a61f509a	3 minutes ago	/bin/sh -c #(nop) WORKDIR /src	0 B
9d9b1f86a696	3 minutes ago	/bin/sh -c #(nop) ENV SERVICE_3000_NAME=self	0 B
2e8721f40082	3 weeks ago	/bin/sh -c apk add --no-cache curl make gcc g	41.44 MB
<missing>	3 weeks ago	/bin/sh -c #(nop) ENV VERSION=v6.2.1 NPM VERS	0 B
<missing>	3 weeks ago	/bin/sh -c #(nop) ADD file:701fd33a2f463fd4bd	4.799 MB

WHY IMAGE ARE SO LIGHTWEIGHT?

- ▶ Each image can share layers
- ▶ An image can be build on top of other images
 - ▶ A static website can be build on top of nginx image which is build on top of an alpine image
- ▶ Building an image could result on rebuilding only last top layers

WHAT'S AN IMAGE

EXAMPLE

redis:alpine

CMD ["redis-server"]

[9 other layers]

ENV REDIS_VERSION=3.2.1

apk add --no-cache 'su-exec>=0.2' [...]

addgroup -S redis && adduser -S -G [...]

nginx:alpine

CMD ["nginx" "-g" "daemon off;"] 0 B

EXPOSE 443/tcp 80/tcp

[5 other layers]

ENV NGINX_VERSION=1.11.1

MAINTAINER NGINX Docker Maintainers

/bin/sh -c #(nop) ADD file:614a9122187935fccf 4.797 MB

alpine:3.3

WHAT'S A CONTAINER

**IT IS A RUNNING IMAGE
WITH A R/W TOP LAYER**

RUN EVERYWHERE

- ▶ Regardless of kernel version
- ▶ Regardless of host distribution
- ▶ Physical or virtual, cloud or not
- ▶ Container and host architecture should match

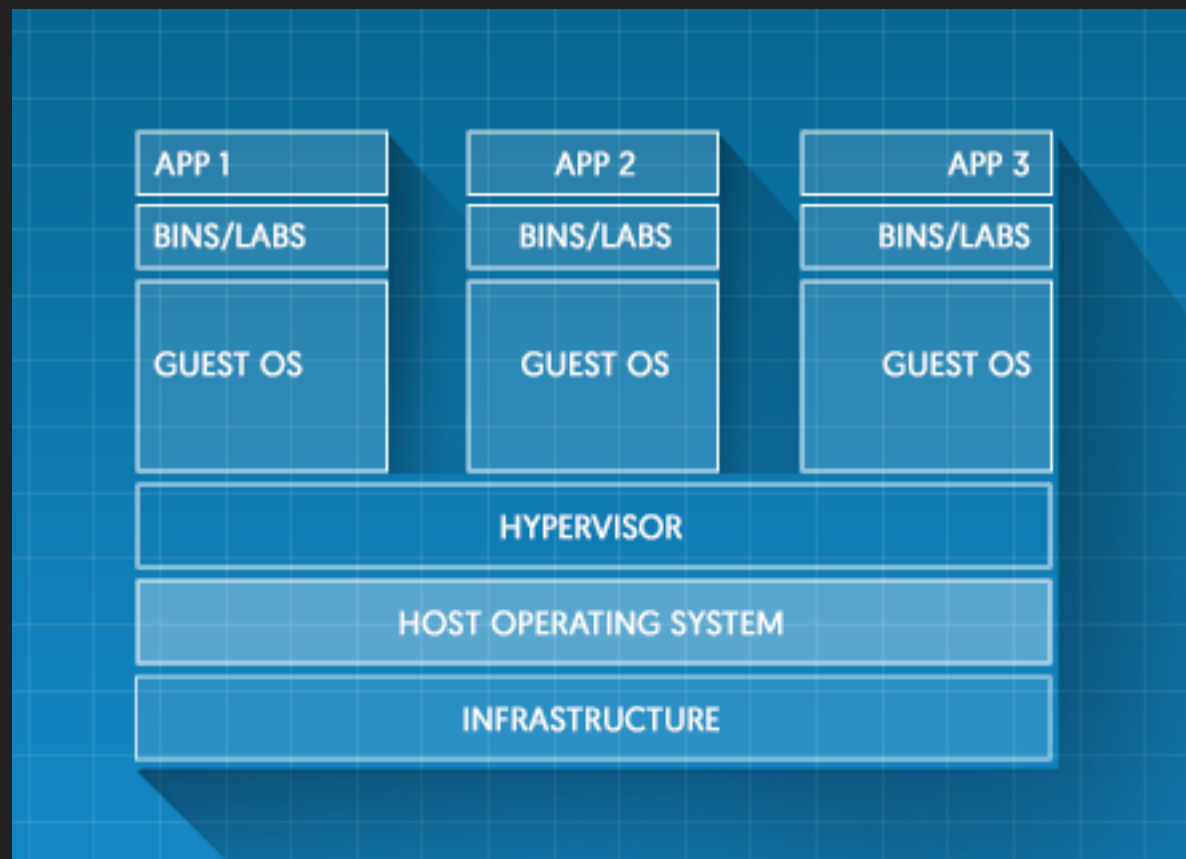
IT'S A LIGHTWEIGHT VM

- ▶ Own process space
- ▶ Own network interface
- ▶ Can run stuff as root or as any custom user

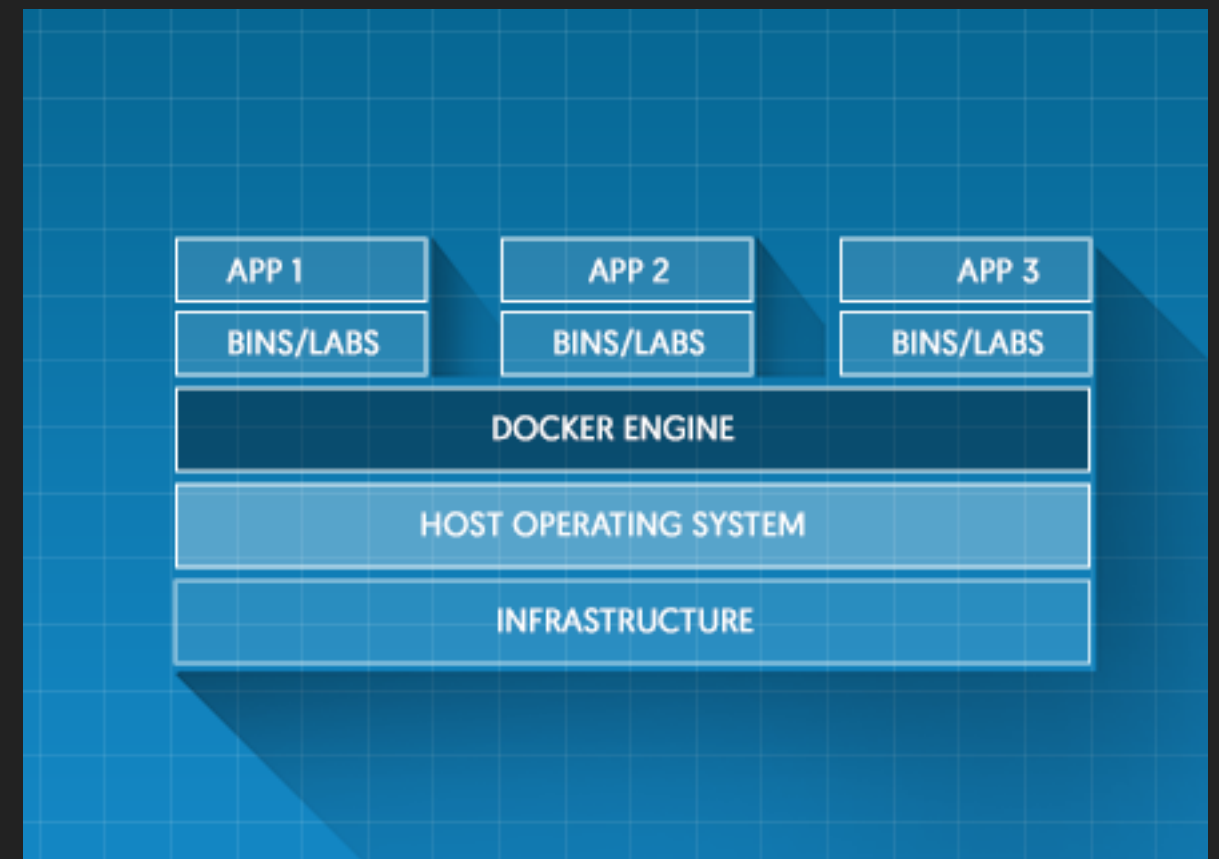
IT'S A CHROOT ON STEROIDS

- ▶ Container = isolated process
- ▶ Share kernel with host (linux and windows containers)
- ▶ No device emulation (neither HVM nor PV)

WHAT'S A CONTAINER



VIRTUAL MACHINES



CONTAINERS

COMPUTE EFFICIENCY

- ▶ CPU performance = native performance
- ▶ Memory performance = ~native performance
- ▶ Network performance = small overhead; can be reduced to zero

STORAGE EFFICIENCY

► Depend on storage driver

AUFS	stable	production-ready	good memory use	smooth Docker experience	high write activity	PaaS-type work
Devicemapper (loop)	stable	in mainline kernel	smooth Docker experience	production	performance	lab testing
Devicemapper (direct-lvm)	stable	production-ready	in mainline kernel	smooth Docker experience	PaaS-type work	
Btrfs	in mainline kernel	high write activity	container churn	build pools		
Overlay	stable	good memory use	in mainline kernel	container churn	lab testing	
ZFS native (ZoL)	PaaS-type work					
ZFS FUSE	stable	lab testing	production			

Key

Has attribute **attribute**

If good for use case **use case**

If bad for use case **use case**

QUESTIONS?